

education

March 16, 2024

Import Data

```
[1]: from ucimlrepo import fetch_ucirepo
import numpy as np
import matplotlib.pyplot as plt
pi = np.pi
import pandas as pd

from ucimlrepo import fetch_ucirepo

import seaborn as sns
%matplotlib inline
```

1 Step 1: Pick (and clean) the data

This is some data about how different social factors impact secondary school performance.

```
[2]: # fetch dataset
student_performance = fetch_ucirepo(id=320)

# data (as pandas dataframes)
X = student_performance.data.features
y = student_performance.data.targets

# metadata
print(student_performance.metadata)

# variable information
print(student_performance.variables)
```

```
{'uci_id': 320, 'name': 'Student Performance', 'repository_url':
'https://archive.ics.uci.edu/dataset/320/student+performance', 'data_url':
'https://archive.ics.uci.edu/static/public/320/data.csv', 'abstract': 'Predict
student performance in secondary education (high school). ', 'area': 'Social
Science', 'tasks': ['Classification', 'Regression'], 'characteristics':
['Multivariate'], 'num_instances': 649, 'num_features': 30, 'feature_types':
['Integer'], 'demographics': ['Sex', 'Age', 'Other', 'Education Level',
```

'Occupation'], 'target_col': ['G1', 'G2', 'G3'], 'index_col': None, 'has_missing_values': 'no', 'missing_values_symbol': None, 'year_of_dataset_creation': 2008, 'last_updated': 'Fri Jan 05 2024', 'dataset_doi': '10.24432/C5TG7T', 'creators': ['Paulo Cortez'], 'intro_paper': {'title': 'Using data mining to predict secondary school student performance', 'authors': 'P. Cortez, A. M. G. Silva', 'published_in': 'Proceedings of 5th Annual Future Business Technology Conference', 'year': 2008, 'url': 'https://www.semanticscholar.org/paper/61d468d5254730bbecf822c6b60d7d6595d9889c', 'doi': None}, 'additional_info': {'summary': 'This data approach student achievement in secondary education of two Portuguese schools. The data attributes include student grades, demographic, social and school related features) and it was collected by using school reports and questionnaires. Two datasets are provided regarding the performance in two distinct subjects: Mathematics (mat) and Portuguese language (por). In [Cortez and Silva, 2008], the two datasets were modeled under binary/five-level classification and regression tasks. Important note: the target attribute G3 has a strong correlation with attributes G2 and G1. This occurs because G3 is the final year grade (issued at the 3rd period), while G1 and G2 correspond to the 1st and 2nd period grades. It is more difficult to predict G3 without G2 and G1, but such prediction is much more useful (see paper source for more details).', 'purpose': None, 'funded_by': None, 'instances_represent': None, 'recommended_data_splits': None, 'sensitive_data': None, 'preprocessing_description': None, 'variable_info': "# Attributes for both student-mat.csv (Math course) and student-por.csv (Portuguese language course) datasets:\r\n1 school - student's school (binary: 'GP' - Gabriel Pereira or 'MS' - Mousinho da Silveira)\r\n2 sex - student's sex (binary: 'F' - female or 'M' - male)\r\n3 age - student's age (numeric: from 15 to 22)\r\n4 address - student's home address type (binary: 'U' - urban or 'R' - rural)\r\n5 famsize - family size (binary: 'LE3' - less or equal to 3 or 'GT3' - greater than 3)\r\n6 Pstatus - parent's cohabitation status (binary: 'T' - living together or 'A' - apart)\r\n7 Medu - mother's education (numeric: 0 - none, 1 - primary education (4th grade), 2 - 5th to 9th grade, 3 - secondary education or 4 - higher education)\r\n8 Fedu - father's education (numeric: 0 - none, 1 - primary education (4th grade), 2 - 5th to 9th grade, 3 - secondary education or 4 - higher education)\r\n9 Mjob - mother's job (nominal: 'teacher', 'health' care related, civil 'services' (e.g. administrative or police), 'at_home' or 'other')\r\n10 Fjob - father's job (nominal: 'teacher', 'health' care related, civil 'services' (e.g. administrative or police), 'at_home' or 'other')\r\n11 reason - reason to choose this school (nominal: close to 'home', school 'reputation', 'course' preference or 'other')\r\n12 guardian - student's guardian (nominal: 'mother', 'father' or 'other')\r\n13 traveltime - home to school travel time (numeric: 1 - <15 min., 2 - 15 to 30 min., 3 - 30 min. to 1 hour, or 4 - >1 hour)\r\n14 studytime - weekly study time (numeric: 1 - <2 hours, 2 - 2 to 5 hours, 3 - 5 to 10 hours, or 4 - >10 hours)\r\n15 failures - number of past class failures (numeric: n if 1<=n<3, else 4)\r\n16 schoolsup - extra educational support (binary: yes or no)\r\n17 famsup - family educational support (binary: yes or no)\r\n18 paid - extra paid classes within the course subject (Math or Portuguese) (binary: yes or no)\r\n19 activities - extra-curricular activities (binary: yes or no)\r\n20 nursery -

attended nursery school (binary: yes or no)\r\n21 higher - wants to take higher education (binary: yes or no)\r\n22 internet - Internet access at home (binary: yes or no)\r\n23 romantic - with a romantic relationship (binary: yes or no)\r\n24 famrel - quality of family relationships (numeric: from 1 - very bad to 5 - excellent)\r\n25 freetime - free time after school (numeric: from 1 - very low to 5 - very high)\r\n26 goout - going out with friends (numeric: from 1 - very low to 5 - very high)\r\n27 Dalc - workday alcohol consumption (numeric: from 1 - very low to 5 - very high)\r\n28 Walc - weekend alcohol consumption (numeric: from 1 - very low to 5 - very high)\r\n29 health - current health status (numeric: from 1 - very bad to 5 - very good)\r\n30 absences - number of school absences (numeric: from 0 to 93)\r\n\r\n# these grades are related with the course subject, Math or Portuguese:\r\n31 G1 - first period grade (numeric: from 0 to 20)\r\n31 G2 - second period grade (numeric: from 0 to 20)\r\n32 G3 - final grade (numeric: from 0 to 20, output target)", 'citation': None}}

	name	role	type	demographic \
0	school	Feature	Categorical	None
1	sex	Feature	Binary	Sex
2	age	Feature	Integer	Age
3	address	Feature	Categorical	None
4	famsize	Feature	Categorical	Other
5	Pstatus	Feature	Categorical	Other
6	Medu	Feature	Integer	Education Level
7	Fedu	Feature	Integer	Education Level
8	Mjob	Feature	Categorical	Occupation
9	Fjob	Feature	Categorical	Occupation
10	reason	Feature	Categorical	None
11	guardian	Feature	Categorical	None
12	traveltime	Feature	Integer	None
13	studytime	Feature	Integer	None
14	failures	Feature	Integer	None
15	schoolsup	Feature	Binary	None
16	famsup	Feature	Binary	None
17	paid	Feature	Binary	None
18	activities	Feature	Binary	None
19	nursery	Feature	Binary	None
20	higher	Feature	Binary	None
21	internet	Feature	Binary	None
22	romantic	Feature	Binary	None
23	famrel	Feature	Integer	None
24	freetime	Feature	Integer	None
25	goout	Feature	Integer	None
26	Dalc	Feature	Integer	None
27	Walc	Feature	Integer	None
28	health	Feature	Integer	None
29	absences	Feature	Integer	None
30	G1	Target	Categorical	None
31	G2	Target	Categorical	None
32	G3	Target	Integer	None

	description	units	missing_values
0	student's school (binary: 'GP' - Gabriel Perei...	None	no
1	student's sex (binary: 'F' - female or 'M' - m...	None	no
2	student's age (numeric: from 15 to 22)	None	no
3	student's home address type (binary: 'U' - urb...	None	no
4	family size (binary: 'LE3' - less or equal to ...	None	no
5	parent's cohabitation status (binary: 'T' - li...	None	no
6	mother's education (numeric: 0 - none, 1 - pr...	None	no
7	father's education (numeric: 0 - none, 1 - pr...	None	no
8	mother's job (nominal: 'teacher', 'health' car...	None	no
9	father's job (nominal: 'teacher', 'health' car...	None	no
10	reason to choose this school (nominal: close t...	None	no
11	student's guardian (nominal: 'mother', 'father...	None	no
12	home to school travel time (numeric: 1 - <15 m...	None	no
13	weekly study time (numeric: 1 - <2 hours, 2 - ...	None	no
14	number of past class failures (numeric: n if 1...	None	no
15	extra educational support (binary: yes or no)	None	no
16	family educational support (binary: yes or no)	None	no
17	extra paid classes within the course subject (...)	None	no
18	extra-curricular activities (binary: yes or no)	None	no
19	attended nursery school (binary: yes or no)	None	no
20	wants to take higher education (binary: yes or...	None	no
21	Internet access at home (binary: yes or no)	None	no
22	with a romantic relationship (binary: yes or no)	None	no
23	quality of family relationships (numeric: from...	None	no
24	free time after school (numeric: from 1 - very...	None	no
25	going out with friends (numeric: from 1 - very...	None	no
26	workday alcohol consumption (numeric: from 1 -...	None	no
27	weekend alcohol consumption (numeric: from 1 -...	None	no
28	current health status (numeric: from 1 - very ...	None	no
29	number of school absences (numeric: from 0 to 93)	None	no
30	first period grade (numeric: from 0 to 20)	None	no
31	second period grade (numeric: from 0 to 20)	None	no
32	final grade (numeric: from 0 to 20, output tar...	None	no

Our data points are different students, and our features for each data point are different characteristics of their family, education, and general parts of their life.

[3]: X

[3]:

	school	sex	age	address	famsize	Pstatus	Medu	Fedu	Mjob	Fjob	\
0	GP	F	18	U	GT3	A	4	4	at_home	teacher	
1	GP	F	17	U	GT3	T	1	1	at_home	other	
2	GP	F	15	U	LE3	T	1	1	at_home	other	
3	GP	F	15	U	GT3	T	4	2	health	services	
4	GP	F	16	U	GT3	T	3	3	other	other	
..			

644	MS	F	19	R	GT3	T	2	3	services	other
645	MS	F	18	U	LE3	T	3	1	teacher	services
646	MS	F	18	U	GT3	T	1	1	other	other
647	MS	M	17	U	LE3	T	3	1	services	services
648	MS	M	18	R	LE3	T	3	2	services	other

	...	higher	internet	romantic	famrel	freetime	goout	Dalc	Walc	health	\
0	...	yes	no	no	4	3	4	1	1	3	
1	...	yes	yes	no	5	3	3	1	1	3	
2	...	yes	yes	no	4	3	2	2	3	3	
3	...	yes	yes	yes	3	2	2	1	1	5	
4	...	yes	no	no	4	3	2	1	2	5	
..	
644	...	yes	yes	no	5	4	2	1	2	5	
645	...	yes	yes	no	4	3	4	1	1	1	
646	...	yes	no	no	1	1	1	1	1	5	
647	...	yes	yes	no	2	4	5	3	4	2	
648	...	yes	yes	no	4	4	1	3	4	5	

	absences
0	4
1	2
2	6
3	0
4	0
..	...
644	4
645	4
646	6
647	6
648	4

[649 rows x 30 columns]

G3 is each individual's final grade in the course, it will be our target data.

[4]: y

	G1	G2	G3
0	0	11	11
1	9	11	11
2	12	13	12
3	14	14	14
4	11	13	13
..
644	10	11	10
645	15	15	16

```
646 11 12 9
647 10 10 10
648 10 11 11
```

```
[649 rows x 3 columns]
```

Want to turn these categorical features into numbers.

```
[ ]: ### Convert objects to categories
obj_columns = X.select_dtypes(['object']).columns

for col in obj_columns:
    X[col] = X[col].astype('category')

### Convert categories to ints
cat_columns = X.select_dtypes(['category']).columns

X[cat_columns] = X[cat_columns].apply(lambda x: x.cat.codes)

# running this cell makes pandas upset at how I'm doing this, so I've hidden
↳ the output
```

I have turned the categorical features, into numerical points in order to perform calculations on them

This could introduce some errors in our conclusions, as, for example, jobs do not necessarily have the same inequality relationship that their numerical representations may encode

(eg. a “veteranarian” may not be greater than or less than a “dentist”, but since they will be transfered to unique integers, one will be encoded with a value greater than or less than the other, additionally, there will be an equally arbitrary linear heirarchy that will arise)

```
[6]: X
```

```
[6]:
```

	school	sex	age	address	famsize	Pstatus	Medu	Fedu	Mjob	Fjob	...	\
0	0	0	18	1	0	0	4	4	0	4
1	0	0	17	1	0	1	1	1	0	2
2	0	0	15	1	1	1	1	1	0	2
3	0	0	15	1	0	1	4	2	1	3
4	0	0	16	1	0	1	3	3	2	2
..
644	1	0	19	0	0	1	2	3	3	2
645	1	0	18	1	1	1	3	1	4	3
646	1	0	18	1	0	1	1	1	2	2
647	1	1	17	1	1	1	3	1	3	3
648	1	1	18	0	1	1	3	2	3	2

	higher	internet	romantic	famrel	freetime	goout	Dalc	Walc	health	\
0	1	0	0	4	3	4	1	1	3	
1	1	1	0	5	3	3	1	1	3	
2	1	1	0	4	3	2	2	3	3	
3	1	1	1	3	2	2	1	1	5	
4	1	0	0	4	3	2	1	2	5	
..	
644	1	1	0	5	4	2	1	2	5	
645	1	1	0	4	3	4	1	1	1	
646	1	0	0	1	1	1	1	1	5	
647	1	1	0	2	4	5	3	4	2	
648	1	1	0	4	4	1	3	4	5	

	absences
0	4
1	2
2	6
3	0
4	0
..	...
644	4
645	4
646	6
647	6
648	4

[649 rows x 30 columns]

[7]: X.dtypes

```
[7]: school      int8
sex            int8
age           int64
address       int8
famsize       int8
Pstatus       int8
Medu         int64
Fedu         int64
Mjob         int8
Fjob         int8
reason       int8
guardian     int8
traveltime   int64
studytime    int64
failures     int64
schoolsup    int8
famsup       int8
```

```

paid            int8
activities      int8
nursery         int8
higher          int8
internet        int8
romantic        int8
famrel          int64
freetime        int64
goout           int64
Dalc            int64
Walc            int64
health          int64
absences        int64
dtype: object

```

Now convert to numpy arrays so we can do math on the data.

```

[8]: Xnp = X.to_numpy()
     print(Xnp)

```

```

[[ 0  0 18 ...  1  3  4]
 [ 0  0 17 ...  1  3  2]
 [ 0  0 15 ...  3  3  6]
 ...
 [ 1  0 18 ...  1  5  6]
 [ 1  1 17 ...  4  2  6]
 [ 1  1 18 ...  4  5  4]]

```

```

[9]: ynp = y.to_numpy()
     print(ynp)

```

```

[[ 0 11 11]
 [ 9 11 11]
 [12 13 12]
 ...
 [11 12  9]
 [10 10 10]
 [10 11 11]]

```

Since Pandas are formatted with different labels/features for each column, and a different sample of data in each row, we transpose the matrix to get it in a form we like better.

```

[10]: Xnp = Xnp.T
      ynp = ynp.T
      print(Xnp.shape)

```

```

(30, 649)

```


2 Part 2: Calculate eigenvalues, of C, histogram them, and fit best M-P distribution

Center the data

```
[11]: Xmean = np.mean(Xnp, axis=1) # the mean of each row/feature across all data
      ↪points
      # changing sape for formatting purposes
      Xmean.shape = (Xmean.size, 1)

      X_circ = Xnp - Xmean
      #print(Xmean.shape)
      print("means of features after centering (should all be approx 0): ", np.
      ↪mean(X_circ, axis=1))
```

```
means of features after centering (should all be approx 0): [-4.37930808e-17
 3.28448106e-17 -1.75172323e-16  4.37930808e-17
 -3.83189457e-17  1.47117381e-17 -1.64224053e-16 -1.53275783e-16
  1.09482702e-16  3.28448106e-17 -3.28448106e-17 -1.16325371e-17
 -6.56896212e-17  1.31379242e-16  4.65301483e-17 -5.47413510e-18
 -5.47413510e-18  1.74488056e-17  3.01077430e-17 -1.23168040e-17
 -6.84266887e-17  5.47413510e-17 -2.46336079e-17 -1.57381384e-16
  1.72435256e-16  8.75861616e-17 -1.25905107e-16  2.73706755e-18
 -1.86120593e-16  8.75861616e-17]
```

Calculate Sample Cov Matrix

```
[12]: m, N = X_circ.shape
      gamma = m/N

      print('number of rows/features, m is:', m)
      print('number of cols/data, N is:', N)
      print('gamma is:', gamma)
      print('1/gamma is:', 1/gamma)

      C = X_circ@X_circ.T/N # Sample Covariance Matrix
      lamb_minus = (1 - np.sqrt(gamma))**2
      lamb_plus = (1 + np.sqrt(gamma))**2

      eig_vals, eig_vecs = np.linalg.eigh(C)

      #print(eig_vals)
      print("The number of eigenvalues is", len(eig_vals))
```

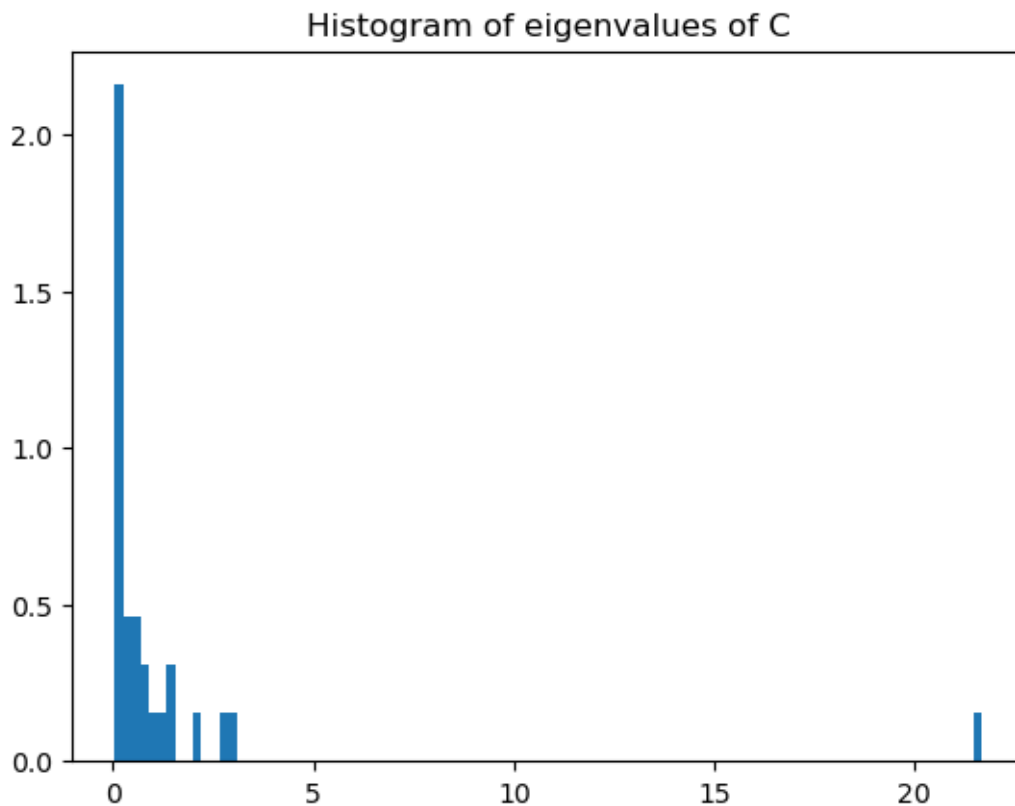
```
number of rows/features, m is: 30
number of cols/data, N is: 649
gamma is: 0.046224961479198766
1/gamma is: 21.633333333333333
The number of eigenvalues is 30
```

```
[13]: print(eig_vals)
```

```
[ 0.05152604  0.0735033  0.08177821  0.08407377  0.12530595  0.13974355
 0.14383113  0.16219291  0.20013945  0.20525621  0.22103834  0.24355781
 0.24670771  0.25531718  0.29479909  0.3718137  0.41579068  0.53536536
 0.61547308  0.67299266  0.76064932  0.84194604  1.05917787  1.33861433
 1.37540904  1.44033446  2.09326526  2.70659398  2.91516739 21.67649624]
```


Histogram of the eigenvalues of C

```
[14]: plt.title("Histogram of eigenvalues of C")
plt.hist(eig_vals,density=True,bins=100, label='Empirical eigenvalues');
```




Histogram of the eigenvalues of C without the big outlier

```
[15]: plt.title("Histogram of eigenvalues of C besides the largest")
plt.hist(eig_vals[:-1],density=True,bins=100, label='Empirical eigenvalues');
```



```
[16]: # Limiting measure
edges = np.linspace(lamb_minus, lamb_plus, 100);
mu = np.sqrt((edges-lamb_minus)*(lamb_plus-edges))/(2*pi*gamma)

# Empirical histogram
plt.figure(figsize=(16,9))
plt.hist(eig_vals, bins=edges, weights=1/(m*(edges[1]-edges[0])*np.ones(m)),
        label='Empirical eigenvalues')
plt.plot(edges, mu/edges, 'r', label='Marcenko-Pastur law')
_ = plt.legend()
```



We can conclude that the data doesn't fit the distribution very well. Very few of the eigenvalues fall within the $\lambda \pm \text{range}$, which is what we expect from Marcenko-Pastur. So likely, our data is not modeled by the “random noise + noisy signal” that we expect, and is mostly the noisy signal part.

3 Part 3: Determining which eigenvalues are outliers

Finding outliers

```
[17]: outliers = [e for e in eig_vals if e < lamb_minus or e > lamb_plus]
print(outliers)
print(len(outliers))
```

```
[0.05152604360203202, 0.07350329636072442, 0.08177820891741958,
0.0840737726825042, 0.12530594637800255, 0.13974355046280207,
0.14383112874409573, 0.16219290931858033, 0.2001394499474628,
0.20525620924943086, 0.2210383426984562, 0.2435578050099582,
0.24670770894905159, 0.25531717628862155, 0.2947990893395774,
0.3718137049165279, 0.4157906798110284, 0.5353653614938254, 0.615473075767388,
2.0932652565320846, 2.706593983342054, 2.9151673949211365, 21.676496239762976]
```

23

Finding outliers that correspond to C_r data from $X = I_m + C_r$

```
[30]: outliers2 = [e for e in outliers if e > np.sqrt(gamma)]
print(outliers2)
print(len(outliers2))
```

```
[0.2210383426984562, 0.2435578050099582, 0.24670770894905159,
0.25531717628862155, 0.2947990893395774, 0.3718137049165279, 0.4157906798110284,
0.5353653614938254, 0.615473075767388, 2.0932652565320846, 2.706593983342054,
2.9151673949211365, 21.676496239762976]
13
```

Can see that 10 of the outlier eigenvalues could have still converged to lambda if 1. we had more data points and 2. the data truly fit the model (we have already concluded that 2. likely already fails, so this is a little pointless)

This means that, according to our model of Noise + Noisy Signal, the top 13 eigenvalues are truly outliers.

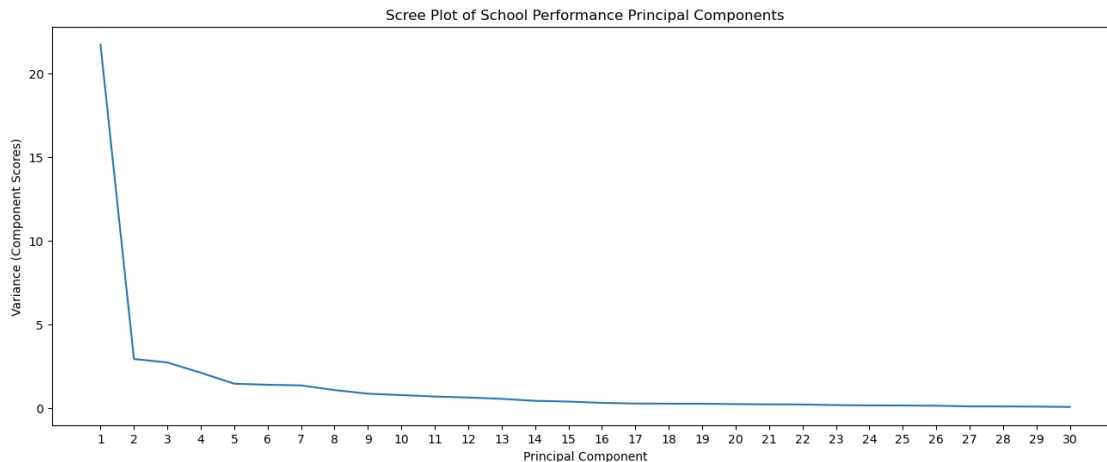
4 Part 4: PCA

```
[18]: U, S, Vt = np.linalg.svd(X_circ, full_matrices=False)
```

```
print(U.shape)
print(S)
print(Vt.shape)
```

```
(30, 30)
[118.60879419  43.49647847  41.91156756  36.85823044  30.57412411
 29.87708931  29.47474686  26.21843693  23.37569201  22.21849245
 20.89909656  19.98604579  18.64006759  16.42705546  15.53406239
 13.83201392  12.87248412  12.65358855  12.57255008  11.97722357
 11.5417191   11.39695148  10.25978548   9.6615942   9.52331687
  9.0179576   7.38673666   7.28519441   6.90678213   5.7827677 ]
(30, 649)
```

```
[19]: plt.figure(figsize=(16, 6))
plt.xticks([i for i in range(1,31)])
plt.xlabel("Principal Component")
plt.ylabel("Variance (Component Scores)")
plt.title("Scree Plot of School Performance Principal Components")
plt.plot([i for i in range(1,31)], np.square(S)/(N-1));
```



We can see visually that there are big(-ish) dropoffs of variance around 2, 5, and 9 PCs, while our choice of outliers says that 13 of them are from the signal.

```
[32]: total_variance = np.sum(np.square(S))/(N-1)

print("total_variance = {:.3f} should approximately equal the sum of feature_
variances: {:.3f}"
      .format(total_variance, np.sum(np.var(Xnp, axis=1))))

two_dim_variance = np.sum(np.square(S[:2]))/(N-1)
three_dim_variance = np.sum(np.square(S[:3]))/(N-1)
five_dim_variance = np.sum(np.square(S[:5]))/(N-1)
nine_dim_variance = np.sum(np.square(S[:9]))/(N-1)
thirteen_dim_variance = np.sum(np.square(S[:13]))/(N-1)

print("The variance of first two components is ", two_dim_variance)
ratio = two_dim_variance / total_variance
print("The ratio two_dim_variance / total_variance = ", ratio)

print("-----")

print("The variance of first three components is ", three_dim_variance)
ratio = three_dim_variance / total_variance
print("The ratio three_dim_variance / total_variance = ", ratio)

print("-----")

print("The variance of first five components is ", five_dim_variance)
ratio = five_dim_variance / total_variance
print("The ratio five_dim_variance / total_variance = ", ratio)
```

```

print("-----")

print("The variance of first nine components is ", nine_dim_variance)
ratio = nine_dim_variance / total_variance
print("The ratio nine_dim_variance / total_variance = ", ratio)

print("-----")

print("The variance of first thirteen components is ", thirteen_dim_variance)
ratio = thirteen_dim_variance / total_variance
print("The ratio thirteen_dim_variance / total_variance = ", ratio)

```

total_variance = 41.412 should approximately equal the sum of feature variances: 41.348

The variance of first two components is 24.62961373288578

The ratio two_dim_variance / total_variance = 0.5947505773272356

The variance of first three components is 27.340384558794725

The ratio three_dim_variance / total_variance = 0.6602096840430871

The variance of first five components is 30.879437362561468

The ratio five_dim_variance / total_variance = 0.7456699645436057

The variance of first nine components is 35.501706776808426

The ratio nine_dim_variance / total_variance = 0.857287525115202

The variance of first thirteen components is 38.090175588449796

The ratio thirteen_dim_variance / total_variance = 0.9197933092827263

As we can see, the first thirteen components account for about 92% of the total variance

```

[34]: fourteen_dim_variance = np.sum(np.square(S[:14]))/(N-1)
fifteen_dim_variance = np.sum(np.square(S[:15]))/(N-1)
sixteen_dim_variance = np.sum(np.square(S[:16]))/(N-1)

print("-----")

print("The variance of first 13 components is ", thirteen_dim_variance)
ratio = thirteen_dim_variance / total_variance
print("The ratio thirteen_dim_variance / total_variance = ", ratio)

print("-----")

print("The variance of first 14 components is ", fourteen_dim_variance)
ratio = fourteen_dim_variance / total_variance
print("The ratio fourteen_dim_variance / total_variance = ", ratio)

```

```

print("-----")

print("The variance of first 15 components is ", fifteen_dim_variance)
ratio = fifteen_dim_variance / total_variance
print("The ratio fifteen_dim_variance / total_variance = ", ratio)

print("-----")

print("The variance of first 16 components is ", sixteen_dim_variance)
ratio = sixteen_dim_variance / total_variance
print("The ratio fifteen_dim_variance / total_variance = ", ratio)

```

```

-----
The variance of first 13 components is  38.090175588449796
The ratio thirteen_dim_variance / total_variance =  0.9197933092827263
-----
The variance of first 14 components is  38.50660792054448
The ratio fourteen_dim_variance / total_variance =  0.9298492270334925
-----
The variance of first 15 components is  38.87899541204267
The ratio fifteen_dim_variance / total_variance =  0.9388415595142194
-----
The variance of first 16 components is  39.1742494382485
The ratio fifteen_dim_variance / total_variance =  0.9459712897831779

```

We can see that the data somewhat supports the conclusion of the first 13 eigenvalues being outliers, since including more PCs (more dimensions) doesn't account for much more of the total variance, only less than 1% each time compared to the much larger increases at previous points (which we continue to not see at later dimensions)

We're going to plot the first 3 Principal Components as a visual aid for the PCA. Trying to observe all 13 dimensions would be unruly to plot.

```

[41]: mu = Xmean
      Q = U[:, [0, 1, 2, 3, 4]]
      print('Q.shape = ', Q.shape)
      print('X_circ.shape = ', X_circ.shape)

      PCs_centered = Q.T @ X_circ

      print('PCs_centered.shape = ', PCs_centered.shape)

      PCs_mean = Q.T @ Xmean
      print('PCs_mean.shape = ', PCs_mean.shape)

      PCs = PCs_mean + PCs_centered
      print('PCs.shape = ', PCs.shape)

```


```
Q.shape = (30, 5)
```



```
X_circ.shape = (30, 649)
PCs_centered.shape = (5, 649)
PCs_mean.shape = (5, 1)
PCs.shape = (5, 649)
```

5 Part 5: Project the Data

```
[22]: plt.figure(figsize=(9, 6))
plt.title("PC2 vs. PC1 for School Performance Data")
plt.xlabel("School Performance PC1")
plt.ylabel("School Performance PC2")
sns.scatterplot(x = PCs[0, :], y = PCs[1, :], hue= y['G3'].to_numpy(),
               palette='inferno');
```



Can see from the above graph, two PCs don't really cut it.

```
[ ]: plt.figure(figsize=(9, 6))
plt.title("PC3 vs. PC1 for School Performance Data")
plt.xlabel("School Performance PC1")
plt.ylabel("School Performance PC3")
sns.scatterplot(x = PCs[0, :], y = PCs[2, :], hue= y['G3'].to_numpy(),
               palette='inferno');
```

```
[ ]: plt.figure(figsize=(9, 6))
plt.title("PC3 vs. PC2 for School Performance Data")
plt.xlabel("School Performance PC2")
plt.ylabel("School Performance PC3")
sns.scatterplot(x = PCs[1, :], y = PCs[2, :], hue= y['G3'].to_numpy(),
↪palette='inferno');
```

```
[25]: PCdf = pd.DataFrame(data=PCs.T, columns=['PC1', 'PC2', 'PC3', 'PC4', 'PC5'])
PCdfplusy = pd.concat([PCdf, y], axis=1)
PCdfplusy
```

```
[25]:
```

	PC1	PC2	PC3	PC4	PC5	G1	G2	G3
0	4.770820	-1.941235	-6.759534	1.095342	-16.863318	0	11	11
1	2.688708	1.005930	-6.492180	0.206895	-15.476193	9	11	11
2	6.729952	0.312722	-6.876292	0.858326	-13.084200	12	13	12
3	0.611806	-2.582423	-5.979268	-1.103219	-14.643675	14	14	14
4	0.698886	-2.674170	-7.190306	-0.957796	-15.126322	11	13	13
..
644	4.785204	-1.829842	-7.938944	-1.188398	-18.217544	10	11	10
645	4.803436	-1.798681	-6.208951	2.492926	-17.499192	15	15	16
646	6.729090	0.412956	-5.819307	-1.920517	-16.827689	11	12	9
647	7.041484	-1.440050	-9.809495	3.073765	-15.445400	10	10	10
648	4.932543	-1.821104	-9.607965	-0.649286	-16.494678	10	11	11


[649 rows x 8 columns]

Here is a 3D scatterplot of the first 3 Principal Components


Even though these 3 PCs account for about 66% of the total variance, which is only about 7% more than the first two, I believe that the visual clarity and split of the data is much more clear.

```
[26]: import plotly.express as px
import plotly.io as pio
pio.renderers.default = "plotly_mimetype+notebook"

fig = px.scatter_3d(PCdfplusy, x='PC1', y='PC2', z='PC3', color='G3')
fig.update_traces(marker_size = 3)
fig.show()
# this isnt showing up when I try to export it to PDF, but I have images of it
↪working below
```



Example Angle:



Another Example Angle:

We can see some clusters of the lower G3 values pulling away from the higher G3 values, however the lowest and highest G3 valued data points still seem to be fairly randomly scattered.

The fifth dimensional data set is hard to visualize, but below is an attempt to view more components that could classify the data further.

```
[27]: num = 5
plt.figure(figsize=(27, 18))
plt.suptitle("Scatter Matrix of PCs")
plt.subplots_adjust(wspace=0.2, hspace=0.3)
for i in range(0, num):
    for j in range(i):
        plt.subplot(num, num, i+num*j)
        sns.scatterplot(y = PCs[i, :], x = PCs[j, :], hue= y['G3'].to_numpy(),
            palette='inferno')
        plt.ylabel("PC" + str(i+1))
        plt.xlabel("PC" + str(j+1))
```

Scatter Matrix of PCs

